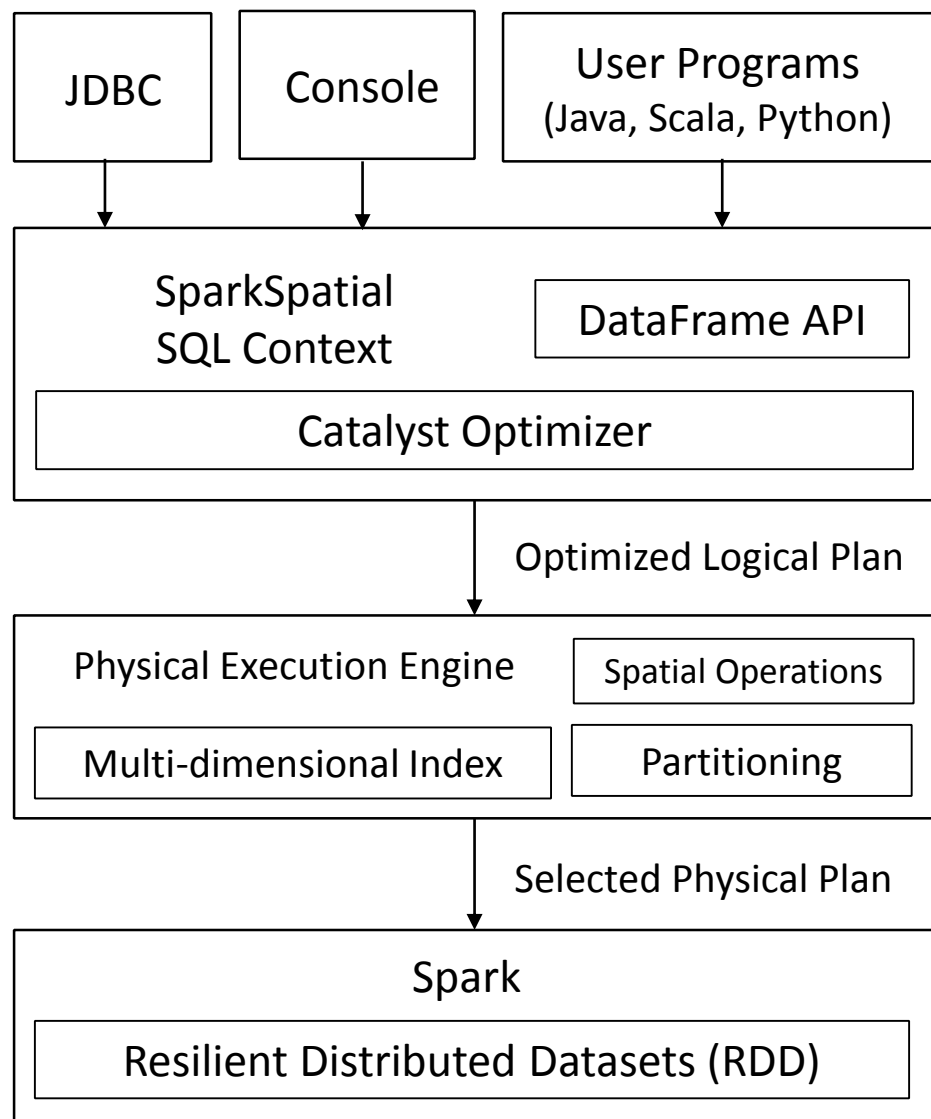


SparkSpatial: Distributed Spatial Data Processing System



- Distributed Spatial Data Processing System based on Spark
- Supporting compound queries on multi-dimensions (spatial, textual, temporal)
- Supporting fast spatial joins on large spatial data
- Adopting a user-friendly SQL-like query language as main interface
- DataFrame API
- Providing several kinds of multi-dimensional index and partitioning strategies
- More complex optimizers based on cost evaluation.

Programming Interface

•SQL-Like Query Language

```
SELECT r.foodtype, count(*) AS c
FROM queries AS q DISTANCE JOIN rests AS r
      ON POINT(r.x, r.y) IN CIRCLERANGE(POINT(q.x, q.y), 10.0)
GROUP BY r.foodtype
ORDER BY c
```

•DataFrame API

```
rest.knn(Point(rest( "x" ), rest( "y" )), Point(3.0, 5.0), 10)
rest.distancejoin(query, Point(query( "x" ), query( "y" )),
                  Point(rest( "x" ), rest( "y" )), 10.0)
    .groupBy(rest( "foodtype" ))
    .select(rest( "foodtype" ), count())
```

•Index Management:

```
CREATE INDEX pointIndex ON point1(x, y) USE rtree
```

Comparison

- Spark SQL:

```
SELECT id, x, y
FROM points
ORDER BY (x - 3) * (x - 3) + (y - 4) * (y - 4)
LIMIT 10
```

- SparkSpatial:

```
SELECT id, x, y
FROM points
WHERE POINT(x, y) in KNN(POINT(3, 4), 10)
```

- Expressing Ability: KNN Join

```
SELECT *
FROM point1 KNN JOIN point2
ON POINT(point2.x, point2.y) IN KNN(POINT(point1.x, point1.y), 10)
```

Indexing Strategy

- Motivation:

Prune useless scanning, save time & resources.

- Challenge:

RDD → Distributed Set → Poor in Random Access

- Solution:

Storage Format Changing & Local + Global Indexing

- Three phases:

- Data Partitioning
- Local Indexing
- Global Indexing

Indexing Strategy (Cont.d)

- Data Partitioning
- Three Main Concerns:
 - Partition Size Fitness: Memory Overflow
 - Data Locality: Query Acceleration.
 - Load Balancing
- Abstract Class: ‘Partitioner’
 - number of partitions + mapping from key to partition
- Solution:
 - One Dimensional: Equal Depth Range Partitioner
 - Multi-Dimensional: STRPartitioner

Indexing Strategy (Cont.d)

- Local Indexing

- Partition Packing:

Pack all elements of a partition into an array.

→ explicit index & faster random access within partition.

- Local Index Building:

Build requested index structure for each packed partition

Co-locate packed partition with its local index.

- As a result:

Storage format for tables changed

RDD[Row] => RDD[PackedPartitionWithIndex]

A Partition => an element with original data & local index

Indexing Strategy (Cont.d)

- Global Indexing

Index for pruning *partitions* in query processing.

- Required Info:

- Partition Boundaries (Data Partitioning)

- Partition Statistics (Local Indexing)

- Index Structure choosing:

- One-dimensional data: sorted range bounds.

- Multi-dimensional data: R-Tree over partition MBRs.

Spatial Operations

- Developer API: PartitionPruningRDD
Skip tasks on specified partition.
- Range Query & Circle Range Query
- Step 1: Query for all partitions that intersect query area.
- Step 2: Invoke range query on remaining partitions.

Spatial Operations (Cont.d)

- KNN Query
- Safe pruning bound: Top k maximum distance.
 - Maximum distance as Distance Function.
 - Step 1: KNN Query on all partition boundaries.
 - Step 2: KNN Query on remaining partition.
 - Take k-th distance as pruning bound.
- Only step 1 is safe, but loose.
- Step 2: Much stronger pruning power:
 - 27 partitions remaining → 4 partitions remaining.

Spatial Operations (Cont.d)

- Distance Join:
R join S on distance between point in R and S less than r.
- Theta-Join => Cartesian Product + Filter (SLOW!!)
- Native Implementation:
 - Nested-Loop Distance Join (with/without R-Tree)
 - SJMR Distance Join
 - R-Tree Distance Join

Spatial Operations (Cont.d)

- KNN Join

R join S on point in S is k-nearest neighbor of point in R over data set S.

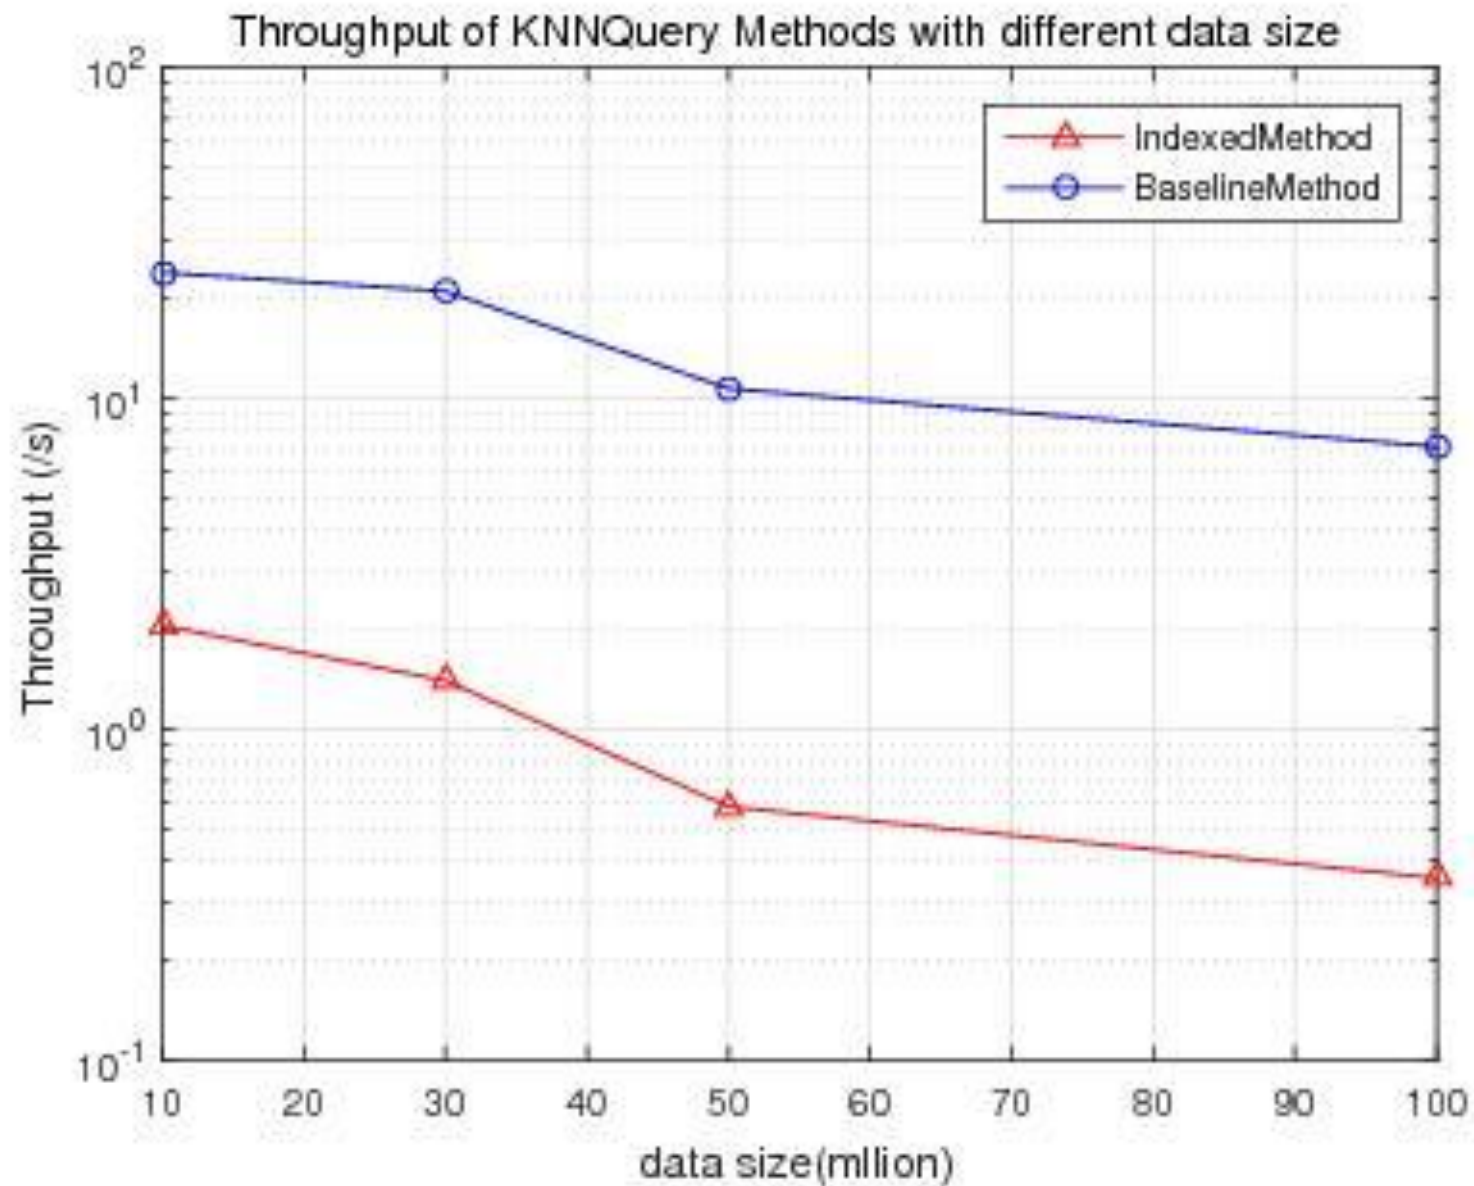
- In Spark SQL: unavailable.

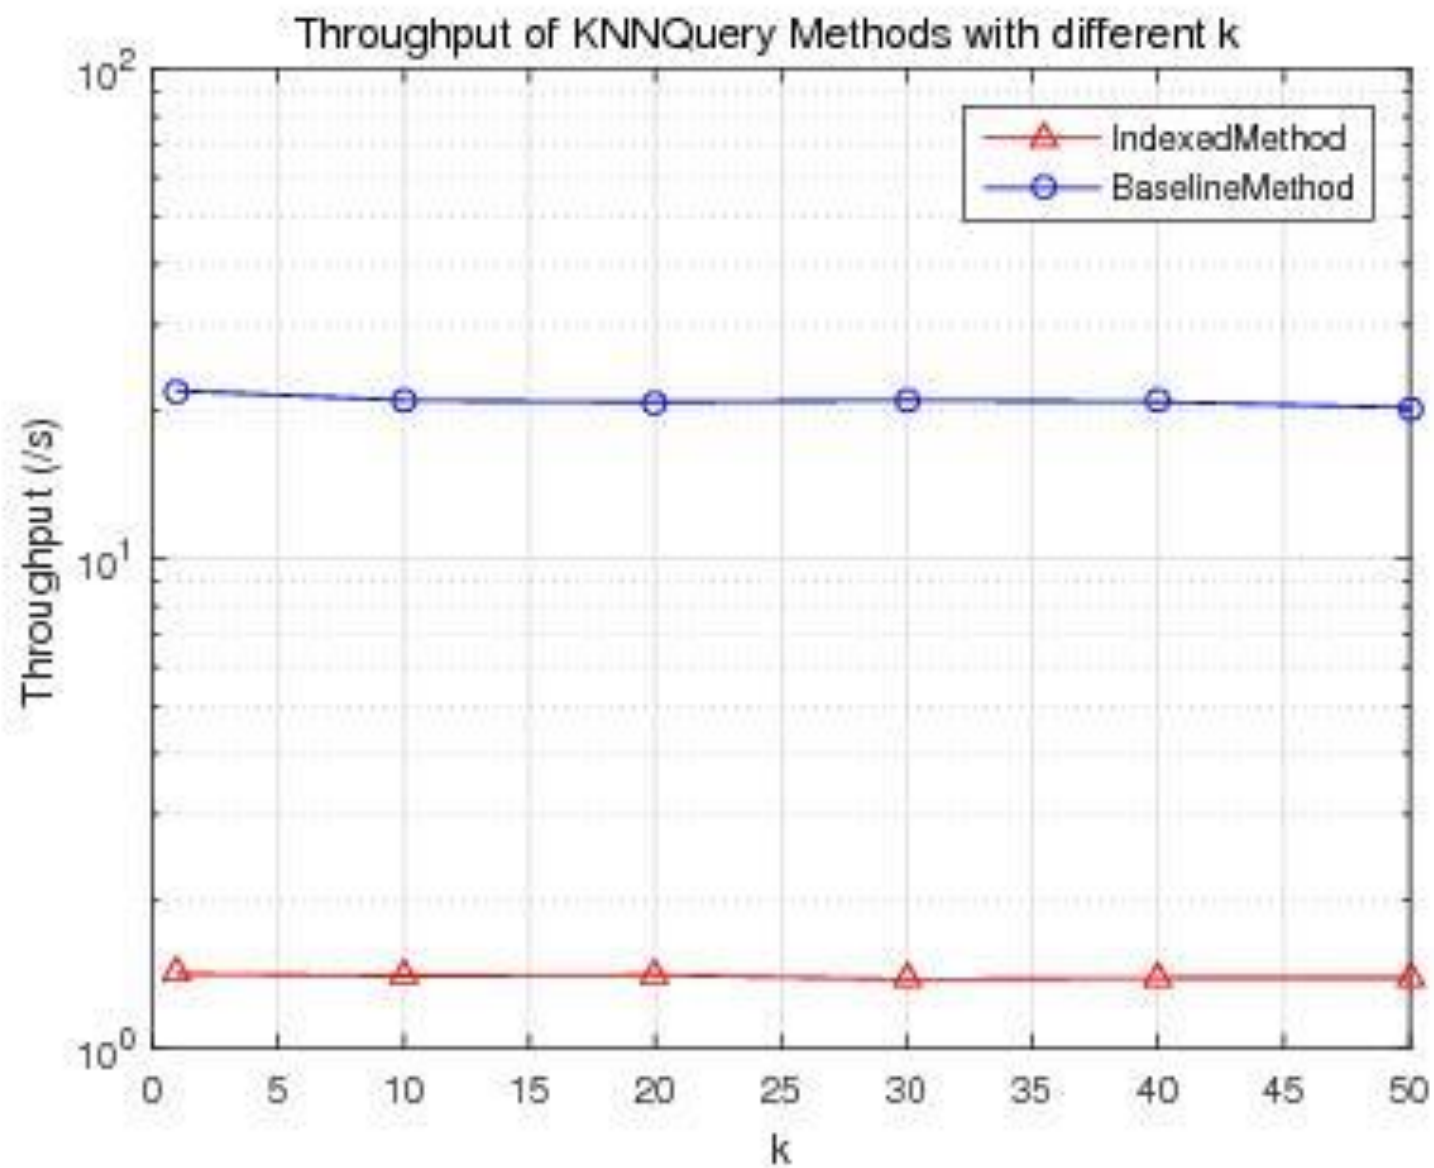
- Implemented Solutions:

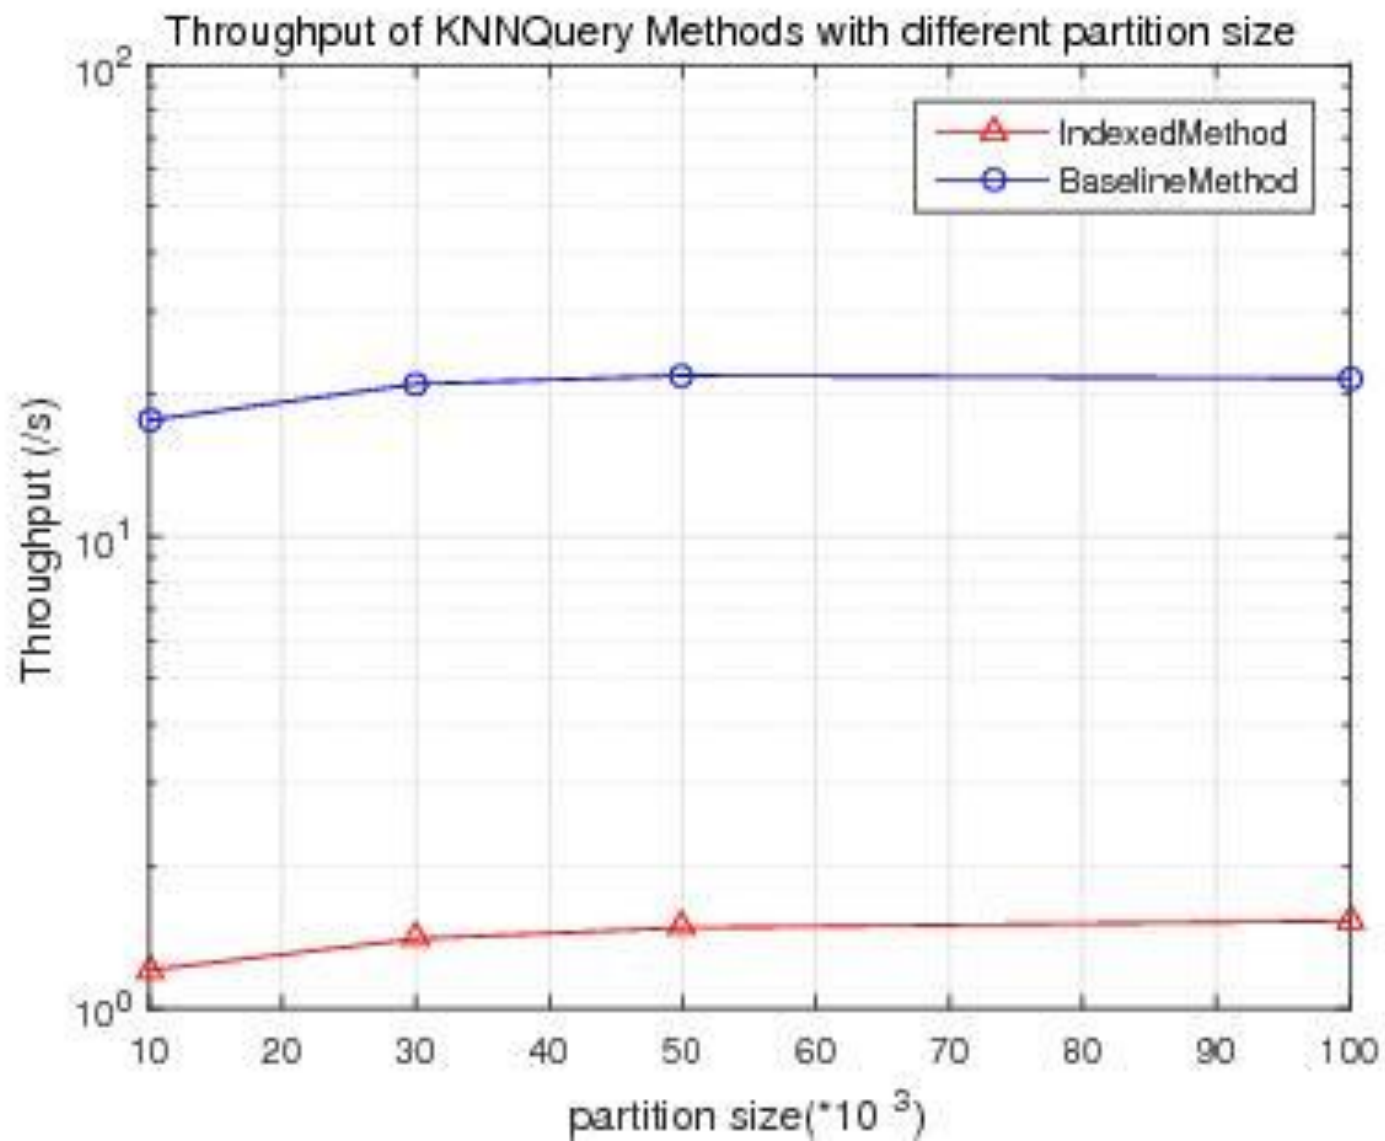
- Cartesian KNN Join
- NestedLoop KNN Join (with/without R-Tree)
- Voronoi KNN Join
- R-Tree KNN Join

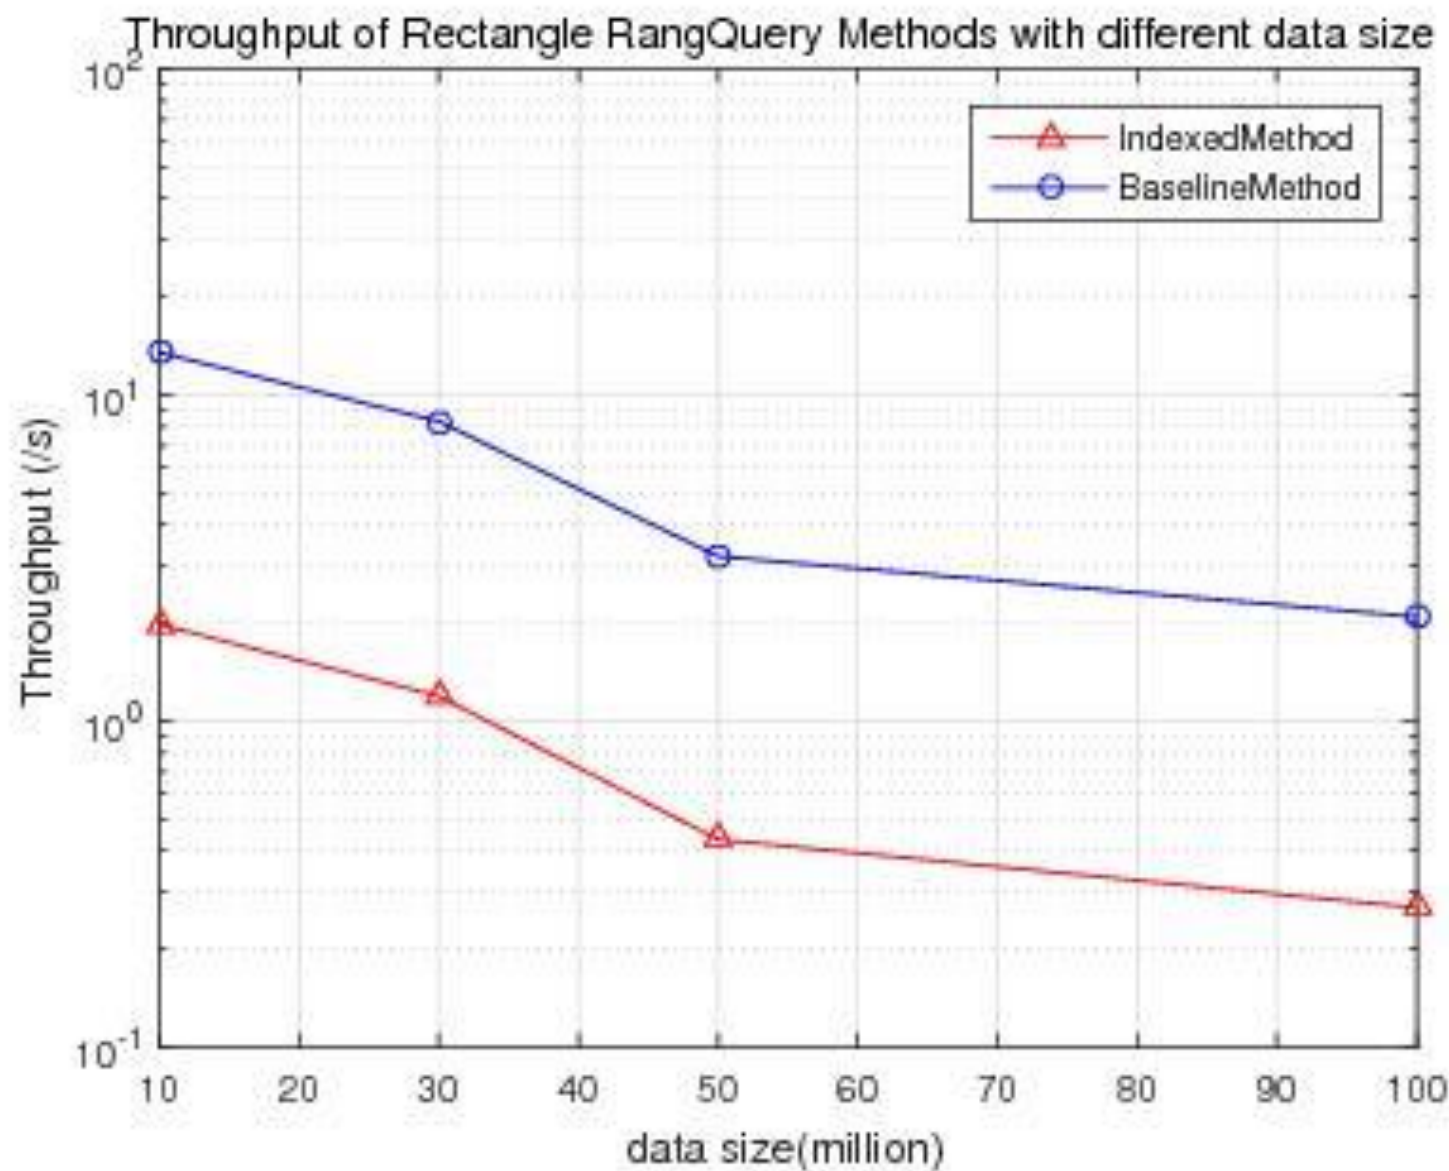
Optimization

- Automatic optimization according to built indexes.
- Formating Query Predicates
 - CNF (Conjunctive Normal Form)
 - DNF (Disconjunctive Normal Form)
- Predicate Combination:
 - e.g. $x > 4 \ \&\& \ x < 6 \ \&\& \ y > 3 \ \&\& \ y < 7$
→ `InRange(Point(4, 3), Point(6, 7))`
- Alternative Execution Path for Indexed Relation:
`IndexedRelationScan`




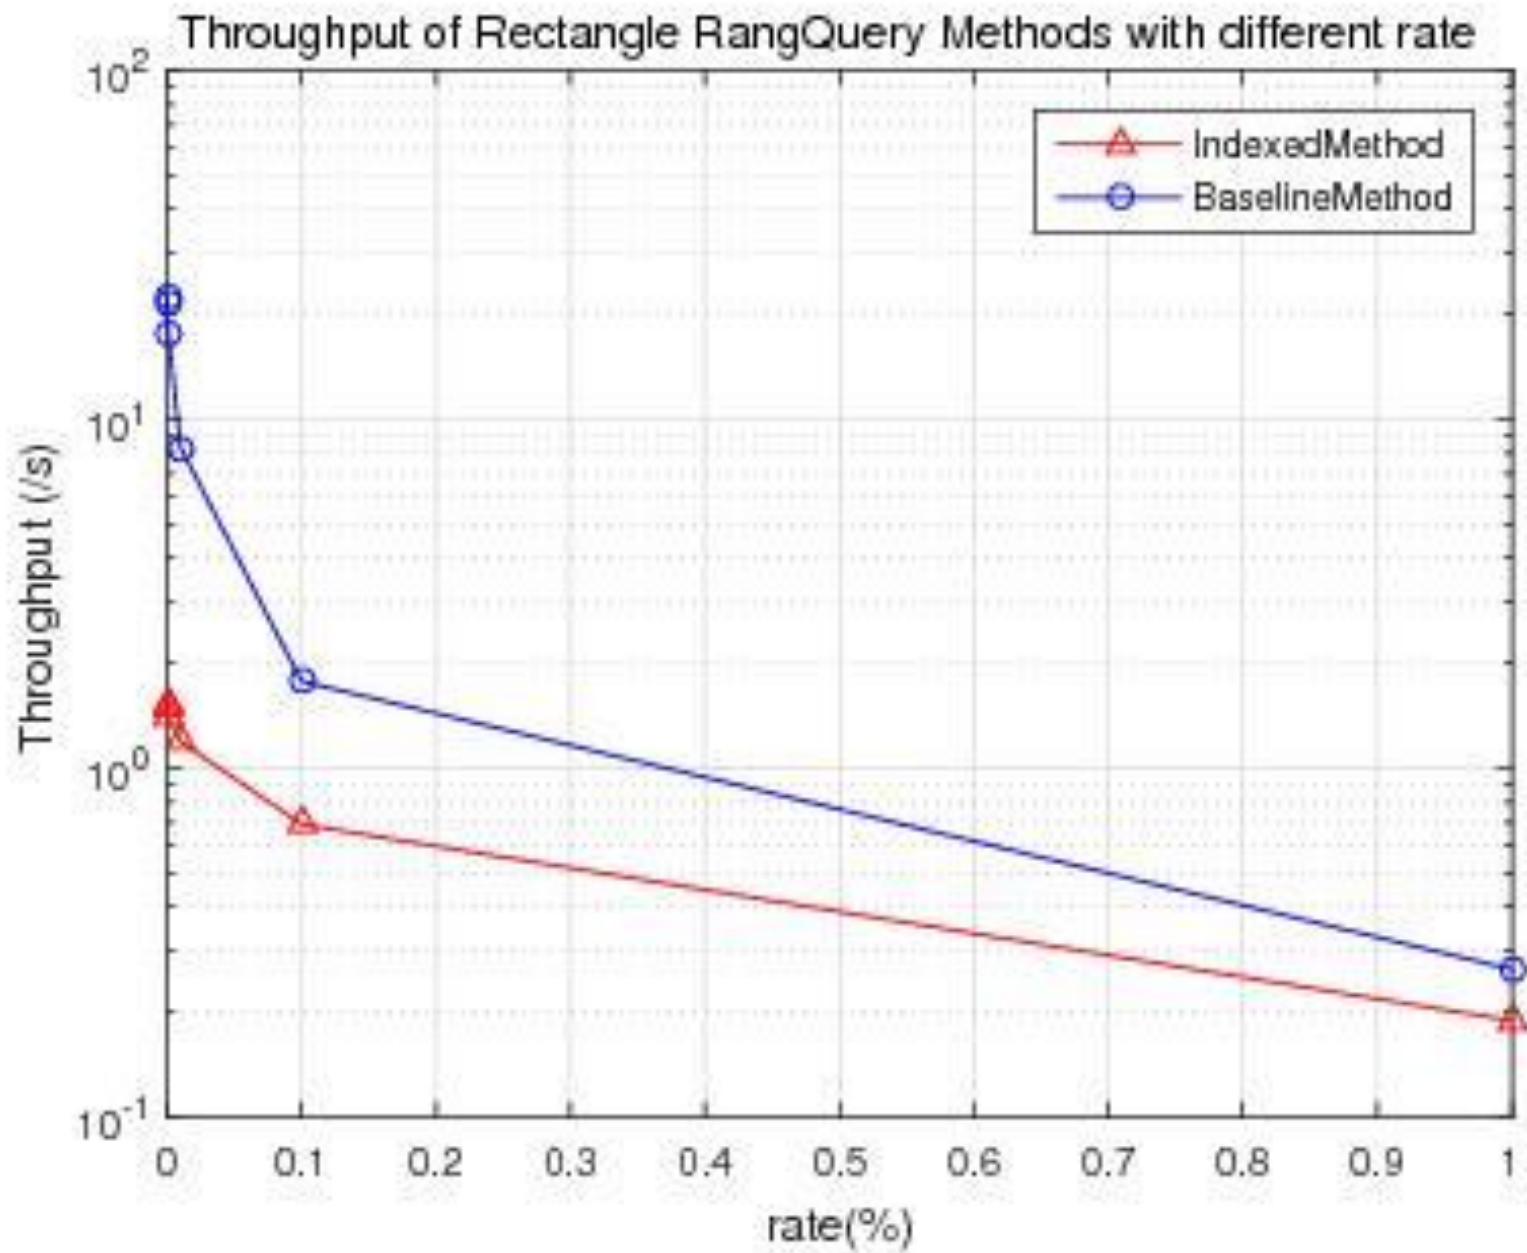




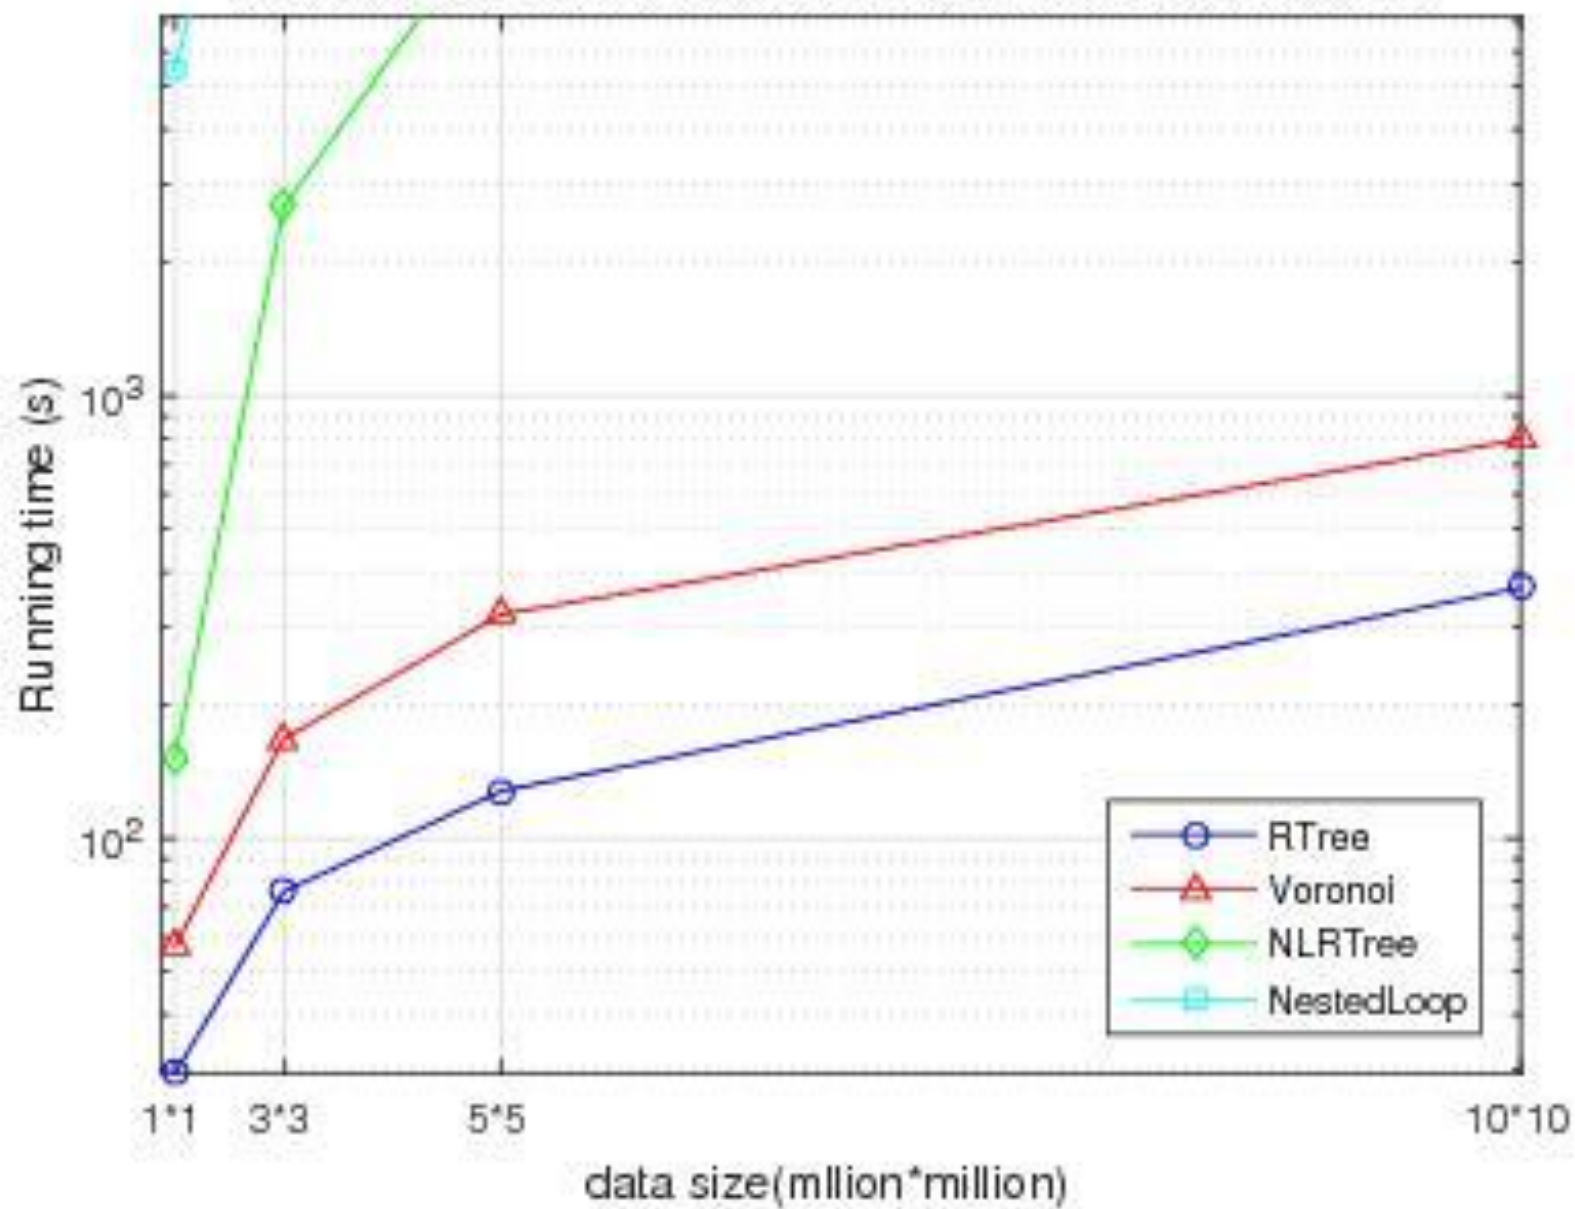


Throughput of Rectangle RangQuery Methods with different partition size







Performance of KNNJoin Methods with different data size



Performance of KNNJoin Methods with different k



Performance of KNNJoin Methods with different partition size



Performance of DistanceJoin Methods with different r

